



COMP417 Lecture 2

Artificial Neural Networks - The Basics

Dr. Hend Dawood


Neural Networks: The Perceptron

Perceptron



Neural Networks: The Perceptron

Perceptron: Simplified model



- Number of inputs combine linearly
 - Threshold logic: Fire if combined input exceeds threshold

$$Y = \begin{cases} 1 & \text{if } \sum_i w_i x_i - T \geq 0 \\ 0 & \text{else} \end{cases}$$

Neural Networks: The Perceptron

Also provided a learning algorithm

$$\mathbf{w} = \mathbf{w} + \eta(d(\mathbf{x}) - y(\mathbf{x}))\mathbf{x}$$


Sequential Learning:

$d(x)$ is the desired output in response to input \mathbf{x}


$y(x)$ is the actual output in response to \mathbf{x}

- Boolean tasks
- Update the weights whenever the perceptron output is wrong
 - Update the weight by the product of the input and the *error* between the desired and actual outputs
- Proved convergence for linearly separable classes


Neural Networks: The Perceptron



Neural Networks: The Perceptron



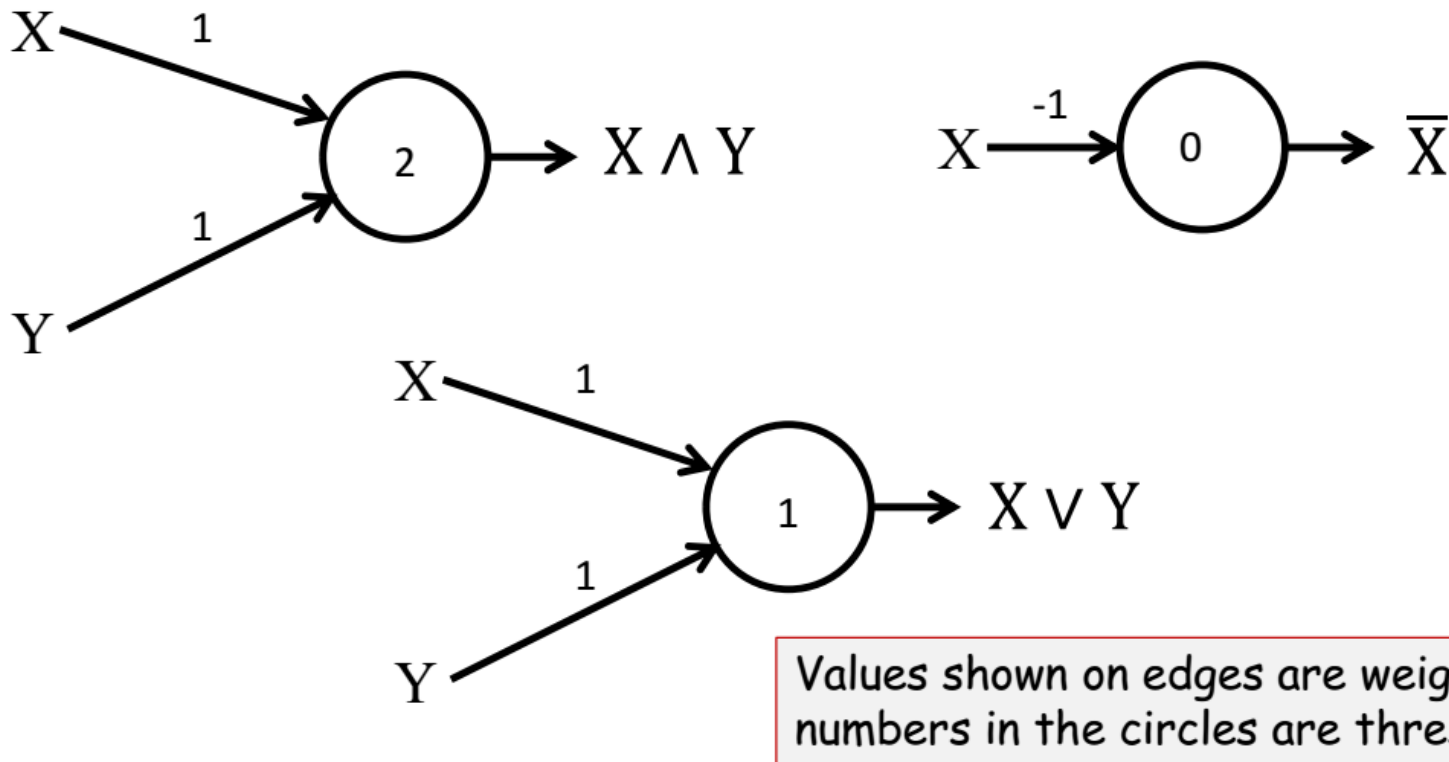
Neural Networks: The Perceptron



AND NOT		
X1	X2	Y
1	1	0
1	0	1
0	1	0
0	0	0

Neural Networks: The Perceptron

Perceptron




- Easily shown to mimic any Boolean gate
- But...

Neural Networks: The Perceptron

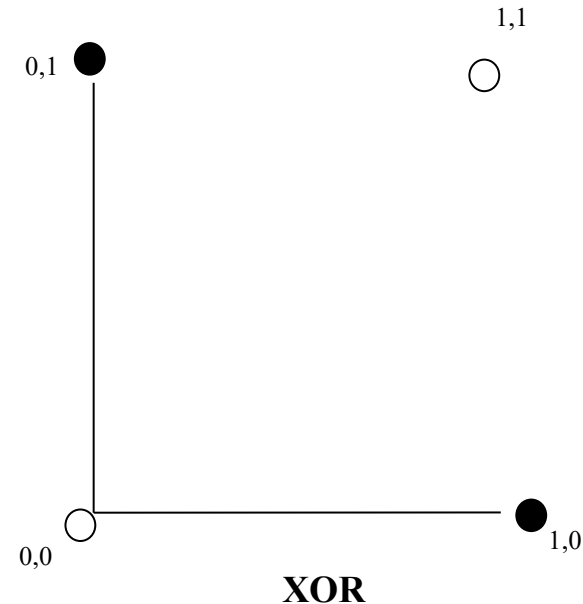

Individual units

No solution for XOR!



Neural Networks: The Perceptron


- What can perceptrons represent?



- Functions which can be separated in this way are called **Linearly Separable**.
- Only linearly Separable functions can be represented by a perceptron.

Neural Networks: The Perceptron



- Example: Training a Perceptron



I_1	I_2	I_3	Summation	Output
-1	0	0	$(-1 * 0.3) + (0 * 0.5) + (0 * -0.4) = -0.3$	0
-1	0	1	$(-1 * 0.3) + (0 * 0.5) + (1 * -0.4) = -0.7$	0
-1	1	0	$(-1 * 0.3) + (1 * 0.5) + (0 * -0.4) = 0.2$	1
-1	1	1	$(-1 * 0.3) + (1 * 0.5) + (1 * -0.4) = -0.2$	0

Neural Networks: The Perceptron


A single neuron is not enough



- Individual elements are weak computational elements
 - Marvin Minsky and Seymour Papert, 1969, *Perceptrons: An Introduction to Computational Geometry*
- *Networked* elements are required

Neural Networks: The Perceptron

Multi-layer Perceptron!




- **XOR**

- The first layer is a “hidden” layer

Neural Networks: The Perceptron

A more generic model


$$((A \& \bar{X} \& Z) | (A \& \bar{Y})) \& ((X \& Y) | (\bar{X} \& \bar{Z}))$$



- A “multi-layer” perceptron
- Can compose arbitrarily complicated Boolean functions!
 - In cognitive terms: Can compute arbitrary Boolean functions over sensory input
 - More on this in the next class

Neural Networks: The Perceptron


But our brain is not Boolean



- We have real inputs
- We make non-Boolean inferences/predictions

Neural Networks: The Perceptron

The perceptron with *real* inputs

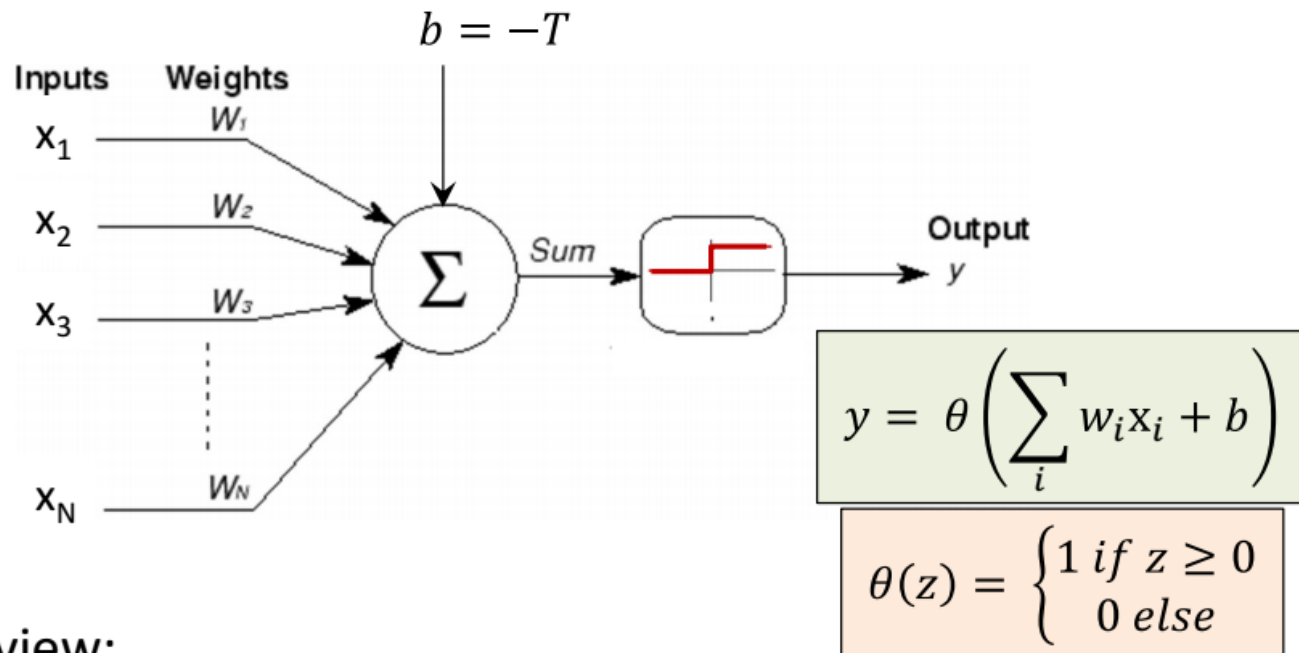


$$y = \begin{cases} 1 & \text{if } \sum_i w_i x_i - T \geq 0 \\ 0 & \text{else} \end{cases}$$

- $x_1 \dots x_N$ are real valued
- $w_1 \dots w_N$ are real valued
- Unit “fires” if weighted input matches (or exceeds) a threshold

Neural Networks: The Perceptron

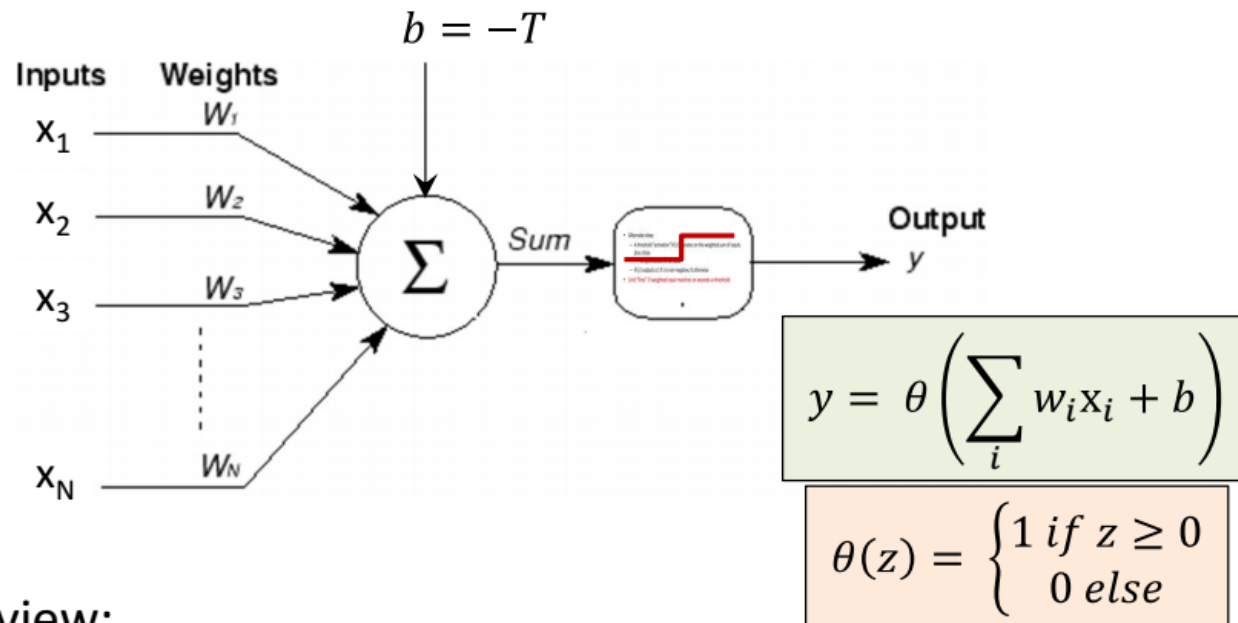
The perceptron with *real* inputs



- Alternate view:
 - A threshold “activation” $\theta(z)$ operates on the weighted sum of inputs plus a bias
 - An *affine* function of the inputs
 - $\theta(z)$ outputs a 1 if z is non-negative, 0 otherwise
- Unit “fires” if weighted input matches or exceeds a threshold

Neural Networks: The Perceptron

The perceptron with *real* inputs






- Alternate view:
 - A threshold “activation” $\theta(z)$ operates on the weighted sum of inputs plus a bias
 - An *affine* function of the inputs
 - $\theta(z)$ outputs a 1 if z is non-negative, 0 otherwise

What is the difference between “linear” and “affine”?

- Unit “fires” if weighted input matches or exceeds a threshold

Neural Networks: The Perceptron

The perceptron with *real* inputs and a *real output*




Activation functions $\sigma(z)$

- $x_1 \dots x_N$ are real valued
- $w_1 \dots w_N$ are real valued
- The output y can also be real valued
- *For now we will continue to assume threshold activations*

Neural Networks: The Perceptron



A Perceptron on Reals



$$y = \begin{cases} 1 & \text{if } \sum_i w_i x_i \geq T \\ 0 & \text{else} \end{cases}$$


- A perceptron operates on *real-valued* vectors

– This is a *linear classifier*



Neural Networks: The Perceptron


Boolean functions with a real perceptron



- Boolean perceptrons are also linear classifiers
 - Purple regions have output 1 in the figures
 - What are these functions
 - Why can we not compose an XOR?

Neural Networks: The Perceptron

Composing complicated “decision” boundaries




Can now be composed into “networks” to compute arbitrary classification “boundaries”

- Build a network of units with a single output that fires if the input is in the coloured area

Neural Networks: The Perceptron


Booleans over the reals



- The network must fire if the input is in the coloured area

Neural Networks: The Perceptron


Booleans over the reals



- The network must fire if the input is in the coloured area

Neural Networks: The Perceptron



Booleans over the reals



- The network must fire if the input is in the coloured area

Neural Networks: The Perceptron



Booleans over the reals



- The network must fire if the input is in the coloured area

Neural Networks: The Perceptron


Booleans over the reals



- The network must fire if the input is in the coloured area

Neural Networks: The Perceptron


Booleans over the reals



- The network must fire if the input is in the coloured area

Neural Networks: The Perceptron


More complex decision boundaries



- Network to fire if the input is in the yellow area
 - “OR” two polygons
 - A third layer is required

Neural Networks: The Perceptron


Complex decision boundaries



- Can compose very complex decision boundaries
 - How complex exactly? More on this in the next class

Neural Networks: The Perceptron


Complex decision boundaries



- Classification problems: finding decision boundaries in high-dimensional space
 - Can be performed by an MLP
- MLPs can *classify* real-valued inputs
- They are universal classifiers
 - For any decision boundary, we can construct an MLP that captures it with arbitrary precision

Neural Networks: The Perceptron


**But what about continuous valued
*outputs?***



- Inputs may be real-valued
- Can outputs be continuous-valued too?

Neural Networks: The Perceptron



MLP as a continuous-valued regression



- A simple 3-unit MLP with a “summing” output unit can generate a “square pulse” over an input
 - Output is 1 only if the input lies between T_1 and T_2
 - T_1 and T_2 can be arbitrarily specified

Neural Networks: The Perceptron

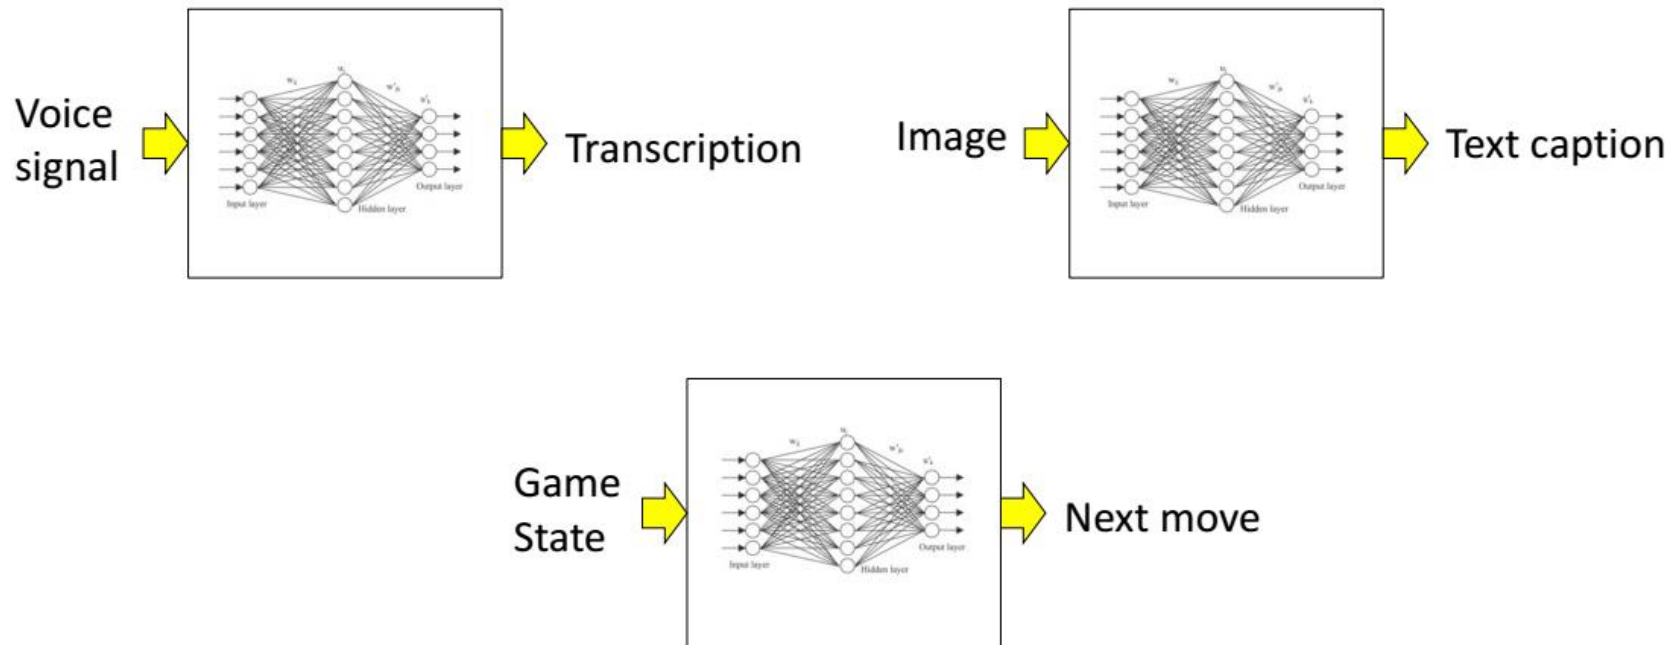
MLP as a continuous-valued regression



- A simple 3-unit MLP can generate a “square pulse” over an input
- **An MLP with many units can model an arbitrary function over an input**
 - To arbitrary precision
 - Simply make the individual pulses narrower
- This generalizes to functions of any number of inputs (next class)

Neural Networks: The Perceptron

These tasks are *functions*



- Each box is actually a function
 - E.g f : Image \rightarrow Caption
 - It can be approximated by a neural network